

XCS251 Assignment 1: Key components of Bitcoin blockchain design

Due Sunday, September 11 at 11:59pm PT.

Guidelines

1. If you have a question about this assignment, we encourage you to post your question on slack channel.
2. Familiarize yourself with the collaboration and honor code policy before starting work.

Submission Instructions

You should submit a PDF with your solutions online in Gradescope. As long as the PDF is legible and organized, the course staff has no preference between a handwritten and a typeset \LaTeX submission. If you wish to typeset your submission and are new to \LaTeX , you can get started with the following:

- Download and install [Tex Live](#) or try [Overleaf](#).
- Submit the compiled PDF.

Honor code

We strongly encourage students to form study groups. Students may discuss and work on assignment in groups. However, each student must write down the solutions independently, and without referring to written notes from the joint session. In other words, each student must understand the solution well enough in order to reconstruct it by him/herself. In addition, each student should write on the assignment the set of people with whom s/he collaborated. Further, because we occasionally reuse problem set questions from previous years, we expect students not to copy, refer to, or look at the solutions in preparing their answers. It is an honor code violation to intentionally refer to a previous year's solutions. More information regarding the Stanford honor code can be found at <https://communitystandards.stanford.edu/policies-and-guidance/honor-code>.

Introduction

In this assignment, there are several questions related to key components of Bitcoin blockchain design. In particular, the first two problems will help you to understand why proof of work hash function needs to be designed carefully and why a binary Merkle tree is good enough to commit a list of elements. The third and fourth question will help you to become familiar with Bitcoin script, ScriptPubKey, ScriptSig, and multisig transactions. The final question examines the minimum proof size required for a lightweight client. By the end of this assignment, you should have a good understanding of proof of work hash function, Merkle proof, Bitcoin script and size of Bitcoin blockchain.

Problem 1. A broken proof of work hash function. In class we discussed using a hash function $H : X \times Y \rightarrow \{0, 1, \dots, 2^n - 1\}$ for a proof of work scheme. Once an $x \in X$ and a difficulty level D are published, it should take an expected D evaluations of the hash function to find a $y \in Y$ such that $H(x, y) < 2^n/D$. Suppose that $X = Y = \{0, 1\}^m$ for some m (say $m = 512$), and consider the hash function

$$H : X \times Y \rightarrow \{0, 1, \dots, 2^{256} - 1\} \quad \text{defined as } H(x, y) := \text{SHA256}(x \oplus y).$$

Here \oplus denotes a bit-wise xor. Show that this H is insecure as a proof of work hash. In particular, suppose D is fixed ahead of time. Show that a clever attacker can find a solution $y \in Y$ with minimal effort once $x \in X$ is published.

Hint: the attacker will do most of the work before x is published.

Problem 2. Beyond binary Merkle trees. Alice can use a binary Merkle tree to commit to a list of elements $S = (T_1, \dots, T_n)$ so that later she can prove to Bob that $S[i] = T_i$ using an inclusion proof containing at most $\lceil \log_2 n \rceil$ hash values. The binding commitment to S is a single hash value. In this question your goal is to explain how to do the same using a k -ary tree, that is, where every non-leaf node has up to k children. The hash value for every non-leaf node is computed as the hash of the concatenation of the values of all its children.

- a. Suppose $S = (T_1, \dots, T_9)$. Explain how Alice computes a commitment to S using a ternary Merkle tree (i.e., $k = 3$). How does Alice later prove to Bob that T_4 is in S ? What values are provided in the proof?
- b. Suppose S contains n elements. What is the length of the proof, i.e., the number of values in the proof, that proves that $S[i] = T_i$, as a function of n and k ? For example, if the proof is $\{x_1, y_2, y_3, y_4\}$ for some Merkle Tree, then the length of the proof is 4.
- c. For large n , if we want to minimize the length of the proof, is it better to use a binary or a ternary tree? Why?

Problem 3. Bitcoin script. Alice is on a backpacking trip and is worried about her devices containing private keys getting stolen. She wants to store her bitcoins in such a way that they can be redeemed via knowledge of a password P . Accordingly, she stores them in the following **ScriptPubKey** address:

```
OP_SHA256
<0xeb271cbcc2340d0b0e6212903e29f22e578ff69bab5690ff654bcd123ae890aa>
OP_EQUAL
```

- a. Write a **ScriptSig** script that will successfully redeem this UTXO given the password P .

Hint: it should only be one line long.

- b. Suppose Alice chooses a six character password P . Explain why her bitcoins can be stolen soon after her UTXO is posted to the blockchain. You may assume that computing SHA256 of all six character passwords can be done in reasonable time.
- c. Suppose Alice chooses a strong 30 character passphrase P . Is the **ScriptPubKey** above a secure way to protect her bitcoins? Why or why not?

Hint: reason through what happens when she tries to redeem her bitcoins.

Problem 4. BitcoinLotto: Suppose the nation of Bitcoinia decides to convert its national lottery to use Bitcoin. A trusted scratch-off ticket printing factory exists and will not keep records of any values printed. Bitcoinia proposes a simple design: a weekly public address is published that holds the jackpot. This allows everyone to verify that the jackpot exists, namely there is a UTXO bound to that address that holds the jackpot amount. Then a weekly run of tickets is printed so that the winning lottery ticket contains the correct private key under the scratch material.

- a. If the winner finds the ticket on Monday and immediately claims the jackpot, this will be bad for sales because players will all realize the lottery has been won. Modify the design to use one of the locktime opcodes to ensure that the prize can only be claimed at the end of the week. (of course, you cannot prevent the winner from proving ownership of the correct private key outside of Bitcoin). The Bitcoin script opcodes are listed here: <https://en.bitcoin.it/wiki/Script>. Alternatively, explain how to use a 2-out-of-2 multisig for this.
- b. Some tickets inevitably get lost or destroyed. Modify the design to roll forward so that any unclaimed jackpot from Week n can be claimed by the winner in Week $n+1$. If the Week $n+1$ jackpot is unclaimed, then the jackpot from both weeks n and $n+1$ can be claimed by the winner of Week $n+2$, and so on. Can you propose a design that works without introducing new ways for the lottery administrators to embezzle funds beyond what is possible on the basic scheme described at the beginning of the

question? You may assume that the lottery system runs for 1000 weeks, and then shuts down permanently.

Hint: use multisig.

Problem 5. Lightweight clients: Suppose Bob runs an ultra lightweight Bitcoin client which receives the current head of the blockchain from a trusted source. This client has very limited memory and so it only stores the block header of the most recent block in the chain (the head of the chain), deleting any previous block headers.

- a. If Alice wants to send a payment to Bob, what information should she include to prove that her payment to Bob has been included in the block chain?
- b. Assume Alice's payment was included in a block k blocks before the current head and there are exactly n transactions per block. (For example, suppose that the head is block number 12, and $k = 12$, then Alice's payment is at block number 0.) Estimate how many bytes this proof will require in terms of n and k . Compute the proof size for $k = 6$ and $n = 1024$.